

```
In [95]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import os
import warnings
warnings.filterwarnings("ignore")
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_auc_score
from sklearn.metrics import accuracy_score
from sklearn.metrics import auc
from sklearn.metrics import plot_roc_curve
from sklearn.metrics import roc_curve
```

```
In [2]: from sklearn.ensemble import AdaBoostClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import StackingClassifier
from sklearn.svm import SVC
from catboost import CatBoostClassifier
from sklearn.neighbors import KNeighborsClassifier
from xgboost import XGBClassifier
from sklearn.linear_model import LogisticRegression
from mlxtend.classifier import EnsembleVoteClassifier
```

```
In [3]: #validation & tuning
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import KFold
from sklearn.model_selection import GridSearchCV

from sklearn.preprocessing import StandardScaler
```

```
In [4]: print(os.getcwd())
os.chdir('D:/OneDrive/BLOGS/League of Legends')
print(os.getcwd())
```

C:\Users\Jinhang Jiang
D:\OneDrive\BLOGS\League of Legends

```
In [373]: info = pd.read_csv("high_diamond_ranked_10min.csv")
```

```
In [295]: data = pd.read_csv("high_diamond_ranked_10min.csv")
```

In [372]: info

Out[372]:

	gameId	blueWins	blueWardsPlaced	blueWardsDestroyed	blueFirstBlood	blueKills	blueDeaths
0	4519157822	0	28	2	1	9	
1	4523371949	0	12	1	0	5	
2	4521474530	0	15	0	0	7	
3	4524384067	0	43	1	0	4	
4	4436033771	0	75	4	0	6	
...
9874	4527873286	1	17	2	1	7	
9875	4527797466	1	54	0	0	6	
9876	4527713716	0	23	1	0	6	
9877	4527628313	0	14	4	1	2	
9878	4523772935	1	18	0	1	6	

9879 rows × 46 columns



In [297]: info.columns

Out[297]: Index(['gameId', 'blueWins', 'blueWardsPlaced', 'blueWardsDestroyed', 'blueFirstBlood', 'blueKills', 'blueDeaths', 'blueAssists', 'blueEliteMonsters', 'blueDragons', 'blueHeralds', 'blueTowersDestroyed', 'blueTotalGold', 'blueAvgLevel', 'blueTotalExperience', 'blueTotalMinionsKilled', 'blueTotalJungleMinionsKilled', 'blueGoldDiff', 'blueExperienceDiff', 'blueCSPerMin', 'blueGoldPerMin', 'redWardsPlaced', 'redWardsDestroyed', 'redFirstBlood', 'redKills', 'redDeaths', 'redAssists', 'redEliteMonsters', 'redDragons', 'redHeralds', 'redTowersDestroyed', 'redTotalGold', 'redAvgLevel', 'redTotalExperience', 'redTotalMinionsKilled', 'redTotalJungleMinionsKilled', 'redGoldDiff', 'redExperienceDiff', 'redCSPerMin', 'redGoldPerMin'], dtype='object')


In [426]: info.blueWins.value_counts(normalize=True)

Out[426]: 0 0.500962
1 0.499038
Name: blueWins, dtype: float64

Plots

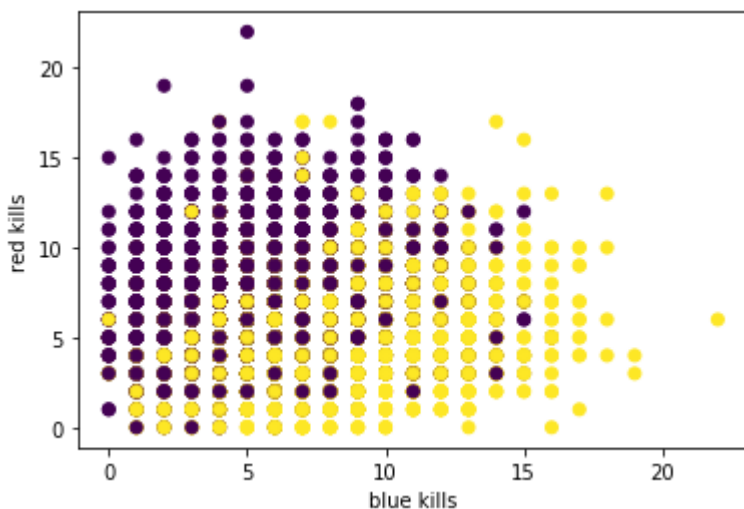
```
In [299]: plt.bar(info.blueWins.unique(), info.blueWins.value_counts(), color=("r", "b"), tick_label=("L", "W"))
```

Out[299]: <BarContainer object of 2 artists>



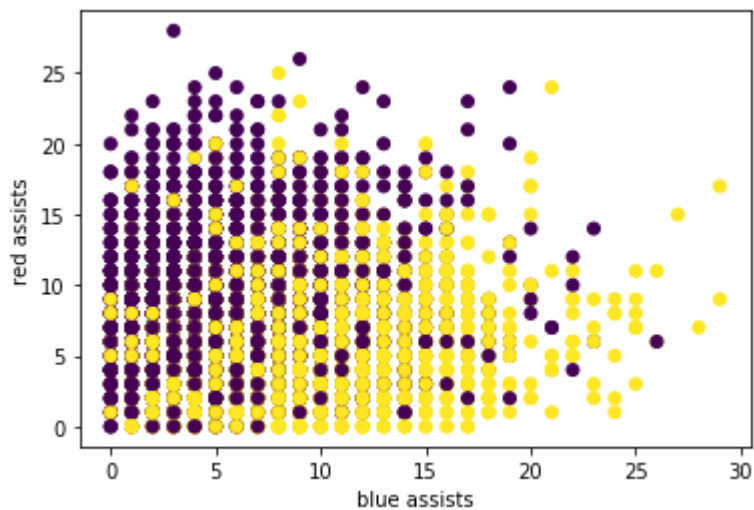
```
In [427]: plt.scatter(info.blueKills, info.redKills, c=info.blueWins)
plt.xlabel("blue kills")
plt.ylabel("red kills")
```

Out[427]: Text(0, 0.5, 'red kills')



```
In [428]: plt.scatter(info.blueAssists, info.redAssists, c=info.blueWins)
plt.xlabel("blue assists")
plt.ylabel("red assists")
```

```
Out[428]: Text(0, 0.5, 'red assists')
```




```
In [302]: info.blueTotalGold.where(info.blueWins==0).sort_values(ascending=False).head(2)
```

```
Out[302]: 9608    21055.0
6137    20887.0
Name: blueTotalGold, dtype: float64
```

```
In [430]: plt.scatter(data.blueTotalGold, data.redTotalGold, c=data.blueWins)
plt.vlines(21056, 12000, 24000, color="red", linestyle=':')
plt.vlines(20000, 12000, 24000, color="red", linestyle=':')
plt.axvspan(21100, 24000, color='g', alpha=0.5, lw=0)
plt.xlabel("blue gold")
plt.ylabel("red gold")
```

Out[430]: Text(0, 0.5, 'red gold')

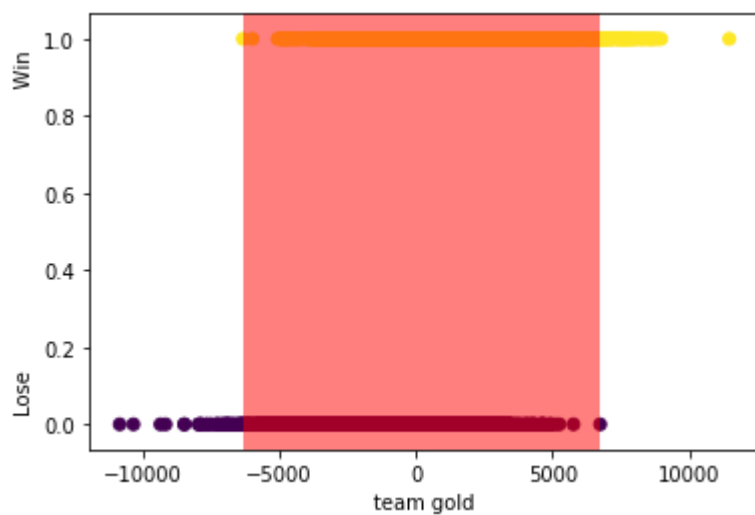


```
In [304]: print(info.blueGoldDiff.where(info.blueWins==0).sort_values(ascending=False).head(1))
print(info.blueGoldDiff.where(info.blueWins==1).sort_values(ascending=True).head(1))
```

```
9608    6744.0
Name: blueGoldDiff, dtype: float64
8459   -6324.0
Name: blueGoldDiff, dtype: float64
```


```
In [434]: plt.scatter(data.blueGoldDiff, data.blueWins, c=data.blueWins)
plt.axvspan(-6324.0, 6744.0, color='r', alpha=0.5, lw=0)
plt.ylabel("Lose" "Win")
plt.xlabel("team gold")
```

Out[434]: Text(0.5, 0, 'team gold')



```
In [439]: plt.scatter(info.blueWardsPlaced, info.redWardsPlaced, c=info.blueWins)
plt.xlabel("blue vision")
plt.ylabel("red vision")
```

Out[439]: Text(0, 0.5, 'red vision')



In [307]: info.columns

```
Out[307]: Index(['gameId', 'blueWins', 'blueWardsPlaced', 'blueWardsDestroyed',
        'blueFirstBlood', 'blueKills', 'blueDeaths', 'blueAssists',
        'blueEliteMonsters', 'blueDragons', 'blueHeralds',
        'blueTowersDestroyed', 'blueTotalGold', 'blueAvgLevel',
        'blueTotalExperience', 'blueTotalMinionsKilled',
        'blueTotalJungleMinionsKilled', 'blueGoldDiff', 'blueExperienceDiff',
        'blueCSPerMin', 'blueGoldPerMin', 'redWardsPlaced', 'redWardsDestroyed',
        'redFirstBlood', 'redKills', 'redDeaths', 'redAssists',
        'redEliteMonsters', 'redDragons', 'redHeralds', 'redTowersDestroyed',
        'redTotalGold', 'redAvgLevel', 'redTotalExperience',
        'redTotalMinionsKilled', 'redTotalJungleMinionsKilled', 'redGoldDiff',
        'redExperienceDiff', 'redCSPerMin', 'redGoldPerMin'],
        dtype='object')
```


```
In [308]: corr_matrix=info.drop(["gameId", 'redWardsPlaced', 'redWardsDestroyed',
        'redFirstBlood', 'redKills', 'redDeaths', 'redAssists',
        'redEliteMonsters', 'redDragons', 'redHeralds', 'redTowersDestroyed',
        'redTotalGold', 'redAvgLevel', 'redTotalExperience',
        'redTotalMinionsKilled', 'redTotalJungleMinionsKilled', 'redGoldDiff',
        'redExperienceDiff', 'redCSPerMin', 'redGoldPerMin'], axis=1).corr(method='pearson')
corr_matrix.iloc[:,0]
```

```
Out[308]: blueWins                1.000000
blueWardsPlaced                0.000087
blueWardsDestroyed             0.044247
blueFirstBlood                 0.201769
blueKills                     0.337358
blueDeaths                    -0.339297
blueAssists                    0.276685
blueEliteMonsters              0.221944
blueDragons                    0.213768
blueHeralds                    0.092385
blueTowersDestroyed            0.115566
blueTotalGold                  0.417213
blueAvgLevel                   0.357820
blueTotalExperience            0.396141
blueTotalMinionsKilled         0.224909
blueTotalJungleMinionsKilled   0.131445
blueGoldDiff                   0.511119
blueExperienceDiff             0.489558
blueCSPerMin                   0.224909
blueGoldPerMin                 0.417213
Name: blueWins, dtype: float64
```

```
In [309]: data.drop(["gameId", 'redWardsPlaced', 'redWardsDestroyed',
                    'redFirstBlood', 'redKills', 'redDeaths', 'redAssists',
                    'redEliteMonsters', 'redDragons', 'redHeralds', 'redTowersDestroyed',
                    'redTotalGold', 'redAvgLevel', 'redTotalExperience',
                    'redTotalMinionsKilled', 'redTotalJungleMinionsKilled', 'redGoldDiff',
                    'redExperienceDiff', 'redCSPerMin', 'redGoldPerMin'], axis=1).describe()
```

Out[309]:

	blueWins	blueWardsPlaced	blueWardsDestroyed	blueFirstBlood	blueKills	blueDeaths
count	9879.000000	9879.000000	9879.000000	9879.000000	9879.000000	9879.000000
mean	0.499038	22.288288	2.824881	0.504808	6.183925	6.137661
std	0.500024	18.019177	2.174998	0.500002	3.011028	2.933811
min	0.000000	5.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	14.000000	1.000000	0.000000	4.000000	4.000000
50%	0.000000	16.000000	3.000000	1.000000	6.000000	6.000000
75%	1.000000	20.000000	4.000000	1.000000	8.000000	8.000000
max	1.000000	250.000000	27.000000	1.000000	22.000000	22.000000



```
In [310]: data.blueWins.where(((data.blueKills+data.blueAssists)/data.blueDeaths)>=3).dropna().value_counts()
```

Out[310]: 1.0 0.765084
0.0 0.234916
Name: blueWins, dtype: float64

```
In [311]: data.blueWins.where(data.blueTowersDestroyed>=1).dropna().value_counts(normalize=True)
```

Out[311]: 1.0 0.75431
0.0 0.24569
Name: blueWins, dtype: float64

```
In [312]: data.blueWins.where(data.blueKills>=8).dropna().value_counts(normalize=True)
```

Out[312]: 1.0 0.69858
0.0 0.30142
Name: blueWins, dtype: float64

Feature Engineering

```
In [388]: info["blueKDA"]=3
for i in range(0, len(info)):
    if data.blueDeaths[i] != 0:
        info["blueKDA"][i]=((data.blueKills[i]+data.blueAssists[i])/data.blueDeaths[i]).round(2)
    else:
        0
```



```
In [390]: info["redKDA"]=3
for i in range(0, len(info)):
    if data.redDeaths[i]!=0:
        info["redKDA"][i] = ((data.redKills[i]+data.redAssists[i])/data.redDeaths[i]).round(2)
    else:
        info["redKDA"][i] = 0
```

```
In [391]: info["blueKDADiff"]=info["blueKDA"]-info["redKDA"]
```

```
In [378]: info["blueGoldAdv"]=3
for i in range(0, len(info)):
    if info["blueTotalGold"][i] >= info["blueTotalGold"][3464]:
        info["blueGoldAdv"][i] = 1
    else:
        info["blueGoldAdv"][i] = 0
```

```
In [379]: info["blueDiffNeg"]=3
for i in range(0, len(info)):
    if info["blueGoldDiff"][i] <= info["blueGoldDiff"][8459]:
        info["blueDiffNeg"][i] = 1
    else:
        info["blueDiffNeg"][i] = 0
```


```
In [380]: info["blueDiffPos"]=3
for i in range(0, 9879):
    if info["blueGoldDiff"][i] >= info["blueGoldDiff"][9608]:
        info["blueDiffPos"][i] = 1
    else:
        info["blueDiffPos"][i] = 0
```

```
In [326]: cols=['gameId', 'blueWins', 'blueWardsPlaced', 'blueWardsDestroyed',
               'blueFirstBlood', 'blueKills', 'blueDeaths', 'blueAssists',
               'blueEliteMonsters', 'blueDragons', 'blueHeralds',
               'blueTowersDestroyed', 'blueTotalGold', 'blueAvgLevel',
               'blueTotalExperience', 'blueTotalMinionsKilled',
               'blueTotalJungleMinionsKilled', 'blueGoldDiff', 'blueExperienceDiff',
               'blueCSPerMin', 'blueGoldPerMin']
info[cols]
```

Out[326]:

	gameId	blueWins	blueWardsPlaced	blueWardsDestroyed	blueFirstBlood	blueKills	blueDeaths
0	4519157822	0	28	2	1	9	
1	4523371949	0	12	1	0	5	
2	4521474530	0	15	0	0	7	
3	4524384067	0	43	1	0	4	
4	4436033771	0	75	4	0	6	
...
9874	4527873286	1	17	2	1	7	
9875	4527797466	1	54	0	0	6	
9876	4527713716	0	23	1	0	6	
9877	4527628313	0	14	4	1	2	
9878	4523772935	1	18	0	1	6	

9879 rows × 21 columns



```
In [392]: cols = ['blueTotalGold', 'blueTotalExperience', 'blueTotalMinionsKilled',
                  'blueGoldDiff', 'blueExperienceDiff', 'blueGoldPerMin',
                  'redTotalGold', 'redTotalExperience', 'redTotalMinionsKilled',
                  'redGoldDiff', 'redExperienceDiff', 'redGoldPerMin']
standard = StandardScaler()
info[cols]=standard.fit_transform(info[cols])
```

```
In [393]: info_x=info.drop(["gameId", "blueWins"], axis=1)
info_y=info.blueWins
```

```
In [329]: #info_x.to_csv("X.csv", index=False)
#info_y.to_csv("Y.csv", index=False)
```

```
In [423]: cv=KFold(n_splits = 5, random_state=2022, shuffle=True)
```

```
In [702]: #train, holdout = train_test_split(info, random_state = 2020, test_size = 0.3)
```

```
In [703]: #train.blueWins.value_counts()
```

```
Out[703]: 0    3509
          1    3406
          Name: blueWins, dtype: int64
```

```
In [704]: #holdout.blueWins.value_counts()
```

```
Out[704]: 1    1524
          0    1440
          Name: blueWins, dtype: int64
```

```
In [705]: #train.to_csv("train.csv", index=False)
          #holdout.to_csv("holdout.csv", index=False)
```

```
In [706]: #y_train = train.blueWins
          #X_train = train.drop(["gameId", "blueWins"], axis=1)
```

```
In [707]: #y_test = holdout.blueWins
          #X_test = holdout.drop(["gameId", "blueWins"], axis=1)
```

KNN

```
In [246]: np.sqrt(len(info_x))
```

```
Out[246]: 99.39315871829409
```

```
In [248]: knn_params = {"n_neighbors": np.arange(95, 105),
                        "weights": ["distance"],
                        "algorithm": ["ball_tree"], #, "kd_tree", "brute"
                        "leaf_size": [1, 2]}
```

```
In [394]: knn = KNeighborsClassifier()
          grid_knn = GridSearchCV(knn, knn_params, cv=5, verbose=2, n_jobs=-1)
          grid_knn.fit(info_x, info_y)
```

Fitting 5 folds for each of 20 candidates, totalling 100 fits

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

[Parallel(n_jobs=-1)]: Done 25 tasks | elapsed: 18.5s

[Parallel(n_jobs=-1)]: Done 100 out of 100 | elapsed: 53.7s finished

```
Out[394]: GridSearchCV(cv=5, estimator=KNeighborsClassifier(), n_jobs=-1,
                    param_grid={'algorithm': ['ball_tree'], 'leaf_size': [1, 2],
                                'n_neighbors': array([ 95,  96,  97,  98,  99, 100, 101, 102,
103, 104]),
                                'weights': ['distance']}},
                    verbose=2)
```

```
In [395]: knn_best_params=grid_knn.best_params_
knn_best_params
```

```
Out[395]: {'algorithm': 'ball_tree',
          'leaf_size': 1,
          'n_neighbors': 98,
          'weights': 'distance'}
```

```
In [396]: knn = KNeighborsClassifier(algorithm="ball_tree",n_neighbors=98,weights="distance",leaf_si
```

```
In [397]: np.mean(cross_val_score(knn,info_x,info_y,cv=cv,scoring="accuracy",n_jobs=-1))
```

```
Out[397]: 0.7061440065597294
```

Logistic Regression

```
In [50]: log_params = {"penalty":["l1","l2","elasticnet"],
                      "solver":["newton-cg", 'lbfgs', 'sag', 'saga'],
                      "C":[1,5,10,20,30,50,100],
                      "warm_start":[True,False]}
```

```
In [51]: logreg=LogisticRegression()
```

```
In [52]: log_grid = GridSearchCV(logreg,log_params,cv=cv,verbose=2,n_jobs=-1)
```

```
In [398]: log_grid.fit(info_x,info_y)
```

Fitting 5 folds for each of 168 candidates, totalling 840 fits

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

[Parallel(n_jobs=-1)]: Done 124 tasks | elapsed: 5.9s

[Parallel(n_jobs=-1)]: Done 840 out of 840 | elapsed: 40.8s finished

```
Out[398]: GridSearchCV(cv=KFold(n_splits=5, random_state=2020, shuffle=False),
                      estimator=LogisticRegression(), n_jobs=-1,
                      param_grid={'C': [1, 5, 10, 20, 30, 50, 100],
                                  'penalty': ['l1', 'l2', 'elasticnet'],
                                  'solver': ['newton-cg', 'lbfgs', 'sag', 'saga'],
                                  'warm_start': [True, False]},
                      verbose=2)
```

```
In [399]: log_grid.best_params_
```

```
Out[399]: {'C': 50, 'penalty': 'l2', 'solver': 'lbfgs', 'warm_start': True}
```

```
In [409]: logreg=LogisticRegression(penalty="l2",solver="sag",C=50, warm_start=False)
```

CatBoostClassifier

```
In [400]: cat_params = {"learning_rate": [0.005],  
                        "depth": [6, 8]}
```

```
In [404]: cat = CatBoostClassifier()
```

```
In [405]: cat_grid = GridSearchCV(cat, cat_params, cv=cv)
```

```
In [406]: cat_grid.fit(info_x, info_y)
```

```
985:   learn: 0.4826309   total: 14.2s   remaining: 202ms  
986:   learn: 0.4826135   total: 14.2s   remaining: 188ms  
987:   learn: 0.4825664   total: 14.3s   remaining: 173ms  
988:   learn: 0.4825577   total: 14.3s   remaining: 159ms  
989:   learn: 0.4825055   total: 14.3s   remaining: 144ms  
990:   learn: 0.4824514   total: 14.3s   remaining: 130ms  
991:   learn: 0.4823349   total: 14.3s   remaining: 115ms  
992:   learn: 0.4823194   total: 14.3s   remaining: 101ms  
993:   learn: 0.4822731   total: 14.3s   remaining: 86.5ms  
994:   learn: 0.4821879   total: 14.4s   remaining: 72.1ms  
995:   learn: 0.4821510   total: 14.4s   remaining: 57.7ms  
996:   learn: 0.4821023   total: 14.4s   remaining: 43.3ms  
997:   learn: 0.4820698   total: 14.4s   remaining: 28.8ms  
998:   learn: 0.4819986   total: 14.4s   remaining: 14.4ms  
999:   learn: 0.4819533   total: 14.4s   remaining: 0ms
```

```
Out[406]: GridSearchCV(cv=KFold(n_splits=5, random_state=2020, shuffle=False),  
                       estimator=<catboost.core.CatBoostClassifier object at 0x0000020F2F2E01C8  
>,  
                       param_grid={'depth': [6, 8], 'learning_rate': [0.005]})
```

```
In [407]: cat_grid.best_params_
```

```
Out[407]: {'depth': 8, 'learning_rate': 0.005}
```

```
In [408]: cat = CatBoostClassifier(learning_rate=0.005, depth=8, verbose=False)
```

SVC

```
In [412]: svc_params={'kernel': ['poly', 'sigmoid'],  
                     'C': [1, 10, 20, 50, 100]}
```

```
In [413]: svc=SVC()
```

```
In [414]: svc_grid=GridSearchCV(svc, svc_params, cv=cv, verbose=2, n_jobs=-1)
```

```
In [415]: svc_grid.fit(info_x, info_y)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

[Parallel(n_jobs=-1)]: Done 25 tasks | elapsed: 47.6s

[Parallel(n_jobs=-1)]: Done 50 out of 50 | elapsed: 1.6min finished

```
Out[415]: GridSearchCV(cv=KFold(n_splits=5, random_state=2020, shuffle=False),
                        estimator=SVC(), n_jobs=-1,
                        param_grid={'C': [1, 10, 20, 50, 100],
                                    'kernel': ['poly', 'sigmoid']},
                        verbose=2)
```

```
In [416]: svc_grid.best_params_
```

```
Out[416]: {'C': 20, 'kernel': 'poly'}
```

```
In [444]: svc=SVC(kernel="poly", probability=True, C=20)
```

```
In [90]: np.mean(cross_val_score(svc, info_x, info_y, cv=cv, scoring="accuracy", n_jobs=-1))
```

```
Out[90]: 0.728008404653308
```

AdaClassifier

```
In [56]: ada_params={"learning_rate": [0.05, 0.1, 0.2, 0.5, 1],
                    "algorithm": ['SAMME', 'SAMME.R'],
                    "n_estimators": [50, 100, 150, 500]}
```

```
In [57]: ada = AdaBoostClassifier()
```

```
In [58]: ada_grid = GridSearchCV(ada, ada_params, cv=cv, verbose=2)
```

```
In [59]: ada_grid.fit(info_x, info_y)

[CV] algorithm=SAMME, learning_rate=0.1, n_estimators=100, total= 1.4s
[CV] algorithm=SAMME, learning_rate=0.1, n_estimators=100 .....
[CV] algorithm=SAMME, learning_rate=0.1, n_estimators=100, total= 1.4s
[CV] algorithm=SAMME, learning_rate=0.1, n_estimators=150 .....
[CV] algorithm=SAMME, learning_rate=0.1, n_estimators=150, total= 2.1s
[CV] algorithm=SAMME, learning_rate=0.1, n_estimators=150 .....
[CV] algorithm=SAMME, learning_rate=0.1, n_estimators=150, total= 2.1s
[CV] algorithm=SAMME, learning_rate=0.1, n_estimators=150 .....
[CV] algorithm=SAMME, learning_rate=0.1, n_estimators=150, total= 2.1s
[CV] algorithm=SAMME, learning_rate=0.1, n_estimators=150 .....
[CV] algorithm=SAMME, learning_rate=0.1, n_estimators=150, total= 2.1s
[CV] algorithm=SAMME, learning_rate=0.1, n_estimators=150 .....
[CV] algorithm=SAMME, learning_rate=0.1, n_estimators=150, total= 2.1s
[CV] algorithm=SAMME, learning_rate=0.1, n_estimators=500 .....
[CV] algorithm=SAMME, learning_rate=0.1, n_estimators=500, total= 7.2s
[CV] algorithm=SAMME, learning_rate=0.1, n_estimators=500 .....
[CV] algorithm=SAMME, learning_rate=0.1, n_estimators=500, total= 7.1s
[CV] algorithm=SAMME, learning_rate=0.1, n_estimators=500 .....
[CV] algorithm=SAMME, learning_rate=0.1, n_estimators=500, total= 7.0s
[CV] algorithm=SAMME, learning_rate=0.1, n_estimators=500
```

```
In [60]: ada_grid.best_params_
```

```
Out[60]: {'algorithm': 'SAMME.R', 'learning_rate': 0.05, 'n_estimators': 150}
```

```
In [445]: ada_best_params = {"learning_rate":0.05,
                             "algorithm":"SAMME.R",
                             "n_estimators":150}
```

```
In [446]: ada=AdaBoostClassifier(**ada_best_params)
```

```
In [67]: np.mean(cross_val_score(ada, info_x, info_y, cv=cv, scoring="accuracy", n_jobs=-1))
```

```
Out[67]: 0.732361041357044
```

Random Forest

```
In [40]: rdf_params={"criterion":["entropy"], #gini
                     "max_depth":[8], #2, 4, 6
                     "min_samples_split":[4], #2, 6
                     "max_features":["sqrt"]} #log2
```

```
In [41]: rdf=RandomForestClassifier()
```

```
In [42]: rdf_grid=GridSearchCV(rdf, rdf_params, cv=cv, verbose=2, n_jobs=-1)
```

```
In [43]: rdf_grid.fit(info_x, info_y)
```

Fitting 5 folds for each of 36 candidates, totalling 180 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 25 tasks      | elapsed:    3.9s
[Parallel(n_jobs=-1)]: Done 146 tasks    | elapsed:   27.0s
[Parallel(n_jobs=-1)]: Done 180 out of 180 | elapsed:   35.9s finished
```

```
Out[43]: GridSearchCV(cv=KFold(n_splits=5, random_state=2020, shuffle=False),
                      estimator=RandomForestClassifier(), n_jobs=-1,
                      param_grid={'criterion': ['gini', 'entropy'],
                                   'max_depth': [4, 6, 8],
                                   'max_features': ['sqrt', 'log2'],
                                   'min_samples_split': [2, 4, 6]},
                      verbose=2)
```

```
In [44]: rdf_grid.best_params_
```

```
Out[44]: {'criterion': 'entropy',
          'max_depth': 8,
          'max_features': 'sqrt',
          'min_samples_split': 4}
```

```
In [447]: rdf=RandomForestClassifier(criterion="entropy", max_depth=8, max_features="sqrt", min_samples
```

```
In [47]: np.mean(cross_val_score(rdf, info_x, info_y, cv=cv, scoring="accuracy", n_jobs=-1))
```

```
Out[47]: 0.7297290524265874
```

XGBClassifier

```
In [68]: xgb_params = {"eta": [0.01, 0.05, 0.1],
                      "max_depth": [2, 4, 6, 8],
                      "objective": ["binary:logistic"],
                      "lambda": [1.20, 50],
                      "alpha": [0, 20],
                      "subsample": [0.3, 0.7, 0.8, 1]}
```

```
In [69]: xgb=XGBClassifier()
```

```
In [70]: xgb_grid=GridSearchCV(xgb, xgb_params, cv=cv, verbose=2, n_jobs=-1)
```



```
In [71]: xgb_grid.fit(info_x, info_y)
```

Fitting 5 folds for each of 192 candidates, totalling 960 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
```

```
[Parallel(n_jobs=-1)]: Done 25 tasks      | elapsed: 11.6s
```

```
[Parallel(n_jobs=-1)]: Done 146 tasks     | elapsed: 1.1min
```

```
[Parallel(n_jobs=-1)]: Done 349 tasks     | elapsed: 2.4min
```

```
[Parallel(n_jobs=-1)]: Done 632 tasks     | elapsed: 4.4min
```

```
[Parallel(n_jobs=-1)]: Done 960 out of 960 | elapsed: 6.6min finished
```

```
Out[71]: GridSearchCV(cv=KFold(n_splits=5, random_state=2020, shuffle=False),
                      estimator=XGBClassifier(base_score=None, booster=None,
                                             colsample_bylevel=None,
                                             colsample_bynode=None,
                                             colsample_bytreet=None, gamma=None,
                                             gpu_id=None, importance_type='gain',
                                             interaction_constraints=None,
                                             learning_rate=None, max_delta_step=None,
                                             max_depth=None, min_child_weight=None,
                                             missing=nan, m...
                                             n_estimators=100, n_jobs=None,
                                             num_parallel_tree=None, random_state=None,
                                             reg_alpha=None, reg_lambda=None,
                                             scale_pos_weight=None, subsample=None,
                                             tree_method=None, validate_parameters=None,
                                             verbosity=None),
                      n_jobs=-1,
                      param_grid={'alpha': [0, 20], 'eta': [0.01, 0.05, 0.1],
                                   'lambda': [1.2, 50], 'max_depth': [2, 4, 6, 8],
                                   'objective': ['binary:logistic'],
                                   'subsample': [0.3, 0.7, 0.8, 1]},
                      verbose=2)
```

```
In [72]: xgb_grid.best_params_
```

```
Out[72]: {'alpha': 0,
          'eta': 0.05,
          'lambda': 50,
          'max_depth': 2,
          'objective': 'binary:logistic',
          'subsample': 1}
```

```
In [448]: xgb=XGBClassifier(eta=0.05, max_depth=2, objective="binary:logistic", reg_lambda=50)
```

Ensemble, Stacking

```
In [424]: eclf = EnsembleVoteClassifier(clfs=[cat, logreg, knn, svc, ada, rdf, xgb], weights=[1, 1, 1, 1, 1, 1],
labels = ['CatBoost', 'Logistic Regression', 'KNN', 'SVC', 'AdaBoost', 'Random Forest', 'XGBoost'])
for clf, label in zip([cat, logreg, knn, svc, ada, rdf, xgb, eclf], labels):

    scores = cross_val_score(clf, info_x, info_y,
                              cv=cv,
                              scoring='accuracy',
                              n_jobs=-1)

    print("[%s] Accuracy: %0.6f (+/- %0.6f) Best: %0.6f "
          % (label, scores.mean(), scores.std(), scores.max()))
```

```
[CatBoost] Accuracy: 0.729832 (+/- 0.010683) Best: 0.739372
[Logistic Regression] Accuracy: 0.730641 (+/- 0.009936) Best: 0.739879
[KNN] Accuracy: 0.706853 (+/- 0.013521) Best: 0.732794
[SVC] Accuracy: 0.721227 (+/- 0.012242) Best: 0.735324
[AdaBoost] Accuracy: 0.732767 (+/- 0.011828) Best: 0.744433
[Random Forest] Accuracy: 0.728313 (+/- 0.007885) Best: 0.738866
[XGBoost] Accuracy: 0.733374 (+/- 0.010998) Best: 0.746964
[Ensemble] Accuracy: 0.730237 (+/- 0.009251) Best: 0.739372
```

```
In [467]: X_train, X_test, y_train, y_test = train_test_split(info_x, info_y, random_state=666, test_si
```

```
In [468]: plt.figure()
lw = 1
#knn
knn.fit(X_train,y_train)
knn_pred = knn.predict_proba(X_test)
fpr, tpr, threshold = roc_curve(y_test,knn_pred[:,1])
roc_auc = auc(fpr, tpr)
plt.plot(fpr, tpr, color='tab:blue',
         lw=lw, label='KNN ROC curve (area = %0.4f)' % roc_auc)

#logreg
logreg.fit(X_train,y_train)
log_pred = logreg.predict_proba(X_test)
fpr, tpr, threshold = roc_curve(y_test,log_pred[:,1])
roc_auc = auc(fpr, tpr)
plt.plot(fpr, tpr, color='tab:green',
         lw=lw, label='Logistic ROC curve (area = %0.4f)' % roc_auc)

#cat
cat.fit(X_train,y_train)
cat_pred = cat.predict_proba(X_test)
fpr, tpr, threshold = roc_curve(y_test,cat_pred[:,1])
roc_auc = auc(fpr, tpr)
plt.plot(fpr, tpr, color='tab:orange',
         lw=lw, label='CatBoost ROC curve (area = %0.4f)' % roc_auc)

#svc
svc.fit(X_train,y_train)
svc_pred = svc.predict_proba(X_test)
fpr, tpr, threshold = roc_curve(y_test,svc_pred[:,1])
roc_auc = auc(fpr, tpr)
plt.plot(fpr, tpr, color='tab:red',
         lw=lw, label='SVC ROC curve (area = %0.4f)' % roc_auc)

#ada
ada.fit(X_train,y_train)
ada_pred = ada.predict_proba(X_test)
fpr, tpr, threshold = roc_curve(y_test,ada_pred[:,1])
roc_auc = auc(fpr, tpr)
plt.plot(fpr, tpr, color='tab:purple',
         lw=lw, label='AdaBoost ROC curve (area = %0.4f)' % roc_auc)

#rdf
rdf.fit(X_train,y_train)
rdf_pred = rdf.predict_proba(X_test)
fpr, tpr, threshold = roc_curve(y_test,rdf_pred[:,1])
roc_auc = auc(fpr, tpr)
plt.plot(fpr, tpr, color='tab:brown',
         lw=lw, label='RandomForest ROC curve (area = %0.4f)' % roc_auc)

#xgb
xgb.fit(X_train,y_train)
xgb_pred = xgb.predict_proba(X_test)
fpr, tpr, threshold = roc_curve(y_test,xgb_pred[:,1])
roc_auc = auc(fpr, tpr)
plt.plot(fpr, tpr, color='tab:gray',
         lw=lw, label='XGBoost ROC curve (area = %0.4f)' % roc_auc)
```

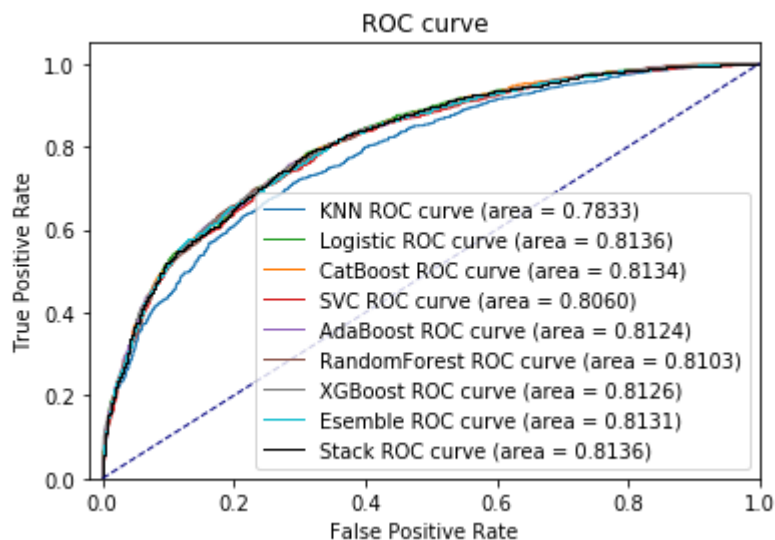
```

#eclf
eclf.fit(X_train,y_train)
eclf_pred = eclf.predict_proba(X_test)
fpr, tpr, threshold = roc_curve(y_test,eclf_pred[:,1])
roc_auc = auc(fpr, tpr)
plt.plot(fpr, tpr, color='tab:cyan',
         lw=lw, label='Esemble ROC curve (area = %0.4f)' % roc_auc)

#stack
stack.fit(X_train,y_train)
stack_pred = stack.predict_proba(X_test)
fpr, tpr, threshold = roc_curve(y_test,stack_pred[:,1])
roc_auc = auc(fpr, tpr)
plt.plot(fpr, tpr, color='black',
         lw=lw, label='Stack ROC curve (area = %0.4f)' % roc_auc)

plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
plt.xlim([-0.02, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC curve')
plt.legend(loc="lower right")
plt.show()

```



```

In [78]: estimators = [
    ('knn', KNeighborsClassifier(algorithm='ball_tree', n_neighbors=98, weights='distance')),
    ('rdf', RandomForestClassifier(criterion="entropy", max_depth=8, max_features="sqrt", min_samples_split=10)),
    ('svc', SVC(kernel="poly", probability=True, C=20)),
    ('cat', CatBoostClassifier(learning_rate=0.005, depth=8, verbose=False)),
    ('ada', AdaBoostClassifier(learning_rate=0.05, algorithm="SAMME.R", n_estimators=150)),
    ('logreg', LogisticRegression(penalty="l2", solver="sag", warm_start=False, C=50)),
    ('xgb', XGBClassifier(eta=0.05, max_depth=2, objective="binary:logistic", reg_lambda=50))
]

stack = StackingClassifier(estimators=estimators,
                          final_estimator=LogisticRegression(penalty="l2", solver="sag", warm_start=False, C=50))

```

```
In [79]: stack_score = cross_val_score(stack, info_x, info_y, cv=cv, scoring='accuracy', n_jobs=-1)
```

```
Out[79]: 0.7321588171987906
```

```
In [89]: print(np.mean(stack_score), np.std(stack_score))  
         print(stack_score.max())
```

```
0.7321588171987906 0.005170871168699806
```

```
0.7388663967611336
```

```

In [440]: tprs = []
          aucs = []
          mean_fpr = np.linspace(0, 1, 100)

          fig, ax = plt.subplots()
          for i, (train, test) in enumerate(cv.split(info_x, info_y)):
              stack.fit(info_x.iloc[train], info_y.iloc[train])
              viz = plot_roc_curve(stack, info_x.iloc[test], info_y.iloc[test],
                                  name='ROC fold {}'.format(i),
                                  alpha=0.3, lw=1, ax=ax)

              interp_tpr = np.interp(mean_fpr, viz.fpr, viz.tpr)
              interp_tpr[0] = 0.0
              tprs.append(interp_tpr)
              aucs.append(viz.roc_auc)


          ax.plot([0, 1], [0, 1], linestyle='--', lw=2, color='r',
                  label='Chance', alpha=.8)

          mean_tpr = np.mean(tprs, axis=0)
          mean_tpr[-1] = 1.0
          mean_auc = auc(mean_fpr, mean_tpr)
          std_auc = np.std(aucs)
          ax.plot(mean_fpr, mean_tpr, color='b',
                  label=r'Mean ROC (AUC = %0.4f  $\pm$  %0.4f)' % (mean_auc, std_auc),
                  lw=2, alpha=.8)

          std_tpr = np.std(tprs, axis=0)
          tprs_upper = np.minimum(mean_tpr + std_tpr, 1)
          tprs_lower = np.maximum(mean_tpr - std_tpr, 0)
          ax.fill_between(mean_fpr, tprs_lower, tprs_upper, color='grey', alpha=.2,
                          label=r' $\pm$  1 std. dev.')

          ax.set(xlim=[-0.05, 1.05], ylim=[-0.05, 1.05],
                  title="Receiver operating characteristic example")
          ax.legend(loc="lower right")
          plt.show()

```



False Prediction Analysis

```
In [469]: stack.fit(X_train,y_train)
          stack_pred = stack.predict_proba(X_test)
```

```
In [470]: X_test["blueWins"]=y_test
          X_test["Prediction"]=stack_pred[:,1].round().astype(int)
```

```
In [471]: X_test["Error"] = 3
          for i in X_test.index:
              if X_test["blueWins"][i] != X_test["Prediction"][i]:
                  X_test["Error"][i] = 1
              else:
                  X_test["Error"][i] = 0
```


```
In [473]: X_test.Error.value_counts()
```

```
Out[473]: 0    2165
          1     799
          Name: Error, dtype: int64
```

```
In [474]: Error = X_test.where(X_test["Error"]==1).dropna()
          Correct = X_test.where(X_test["Error"]==0).dropna()
          Error_normal = data.loc[Error.index]
          Correct_normal = data.loc[Correct.index]
          data_normal = data.loc[X_test.index]
```

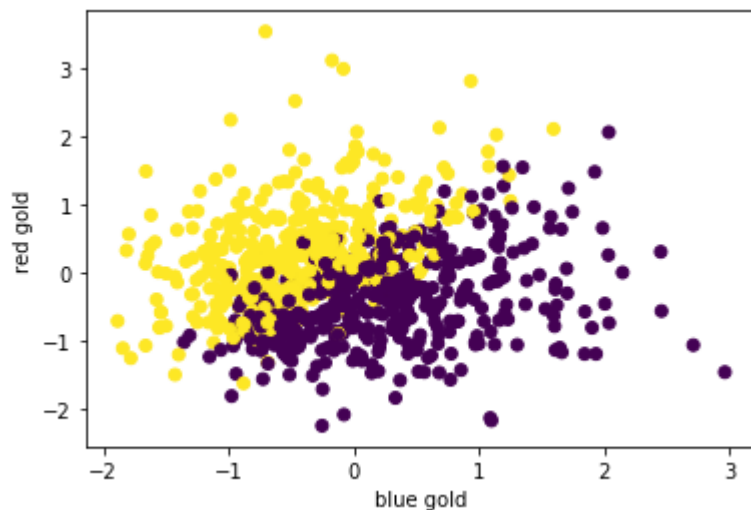
```
In [475]: plt.bar(Error.blueWins.unique(),Error.blueWins.value_counts(),
                  color=("r","b"),tick_label=("L","W"))
          print(Error.blueWins.value_counts(normalize=True))
```

```
0.0    0.519399
1.0    0.480601
          Name: blueWins, dtype: float64
```



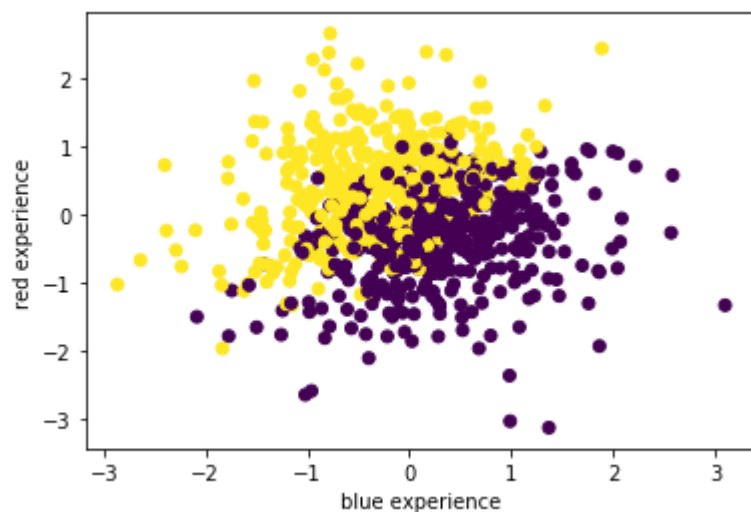
```
In [476]: plt.scatter(Error.blueTotalGold, Error.redTotalGold, c=Error.blueWins)
plt.xlabel("blue gold")
plt.ylabel("red gold")
```

```
Out[476]: Text(0, 0.5, 'red gold')
```



```
In [477]: plt.scatter(Error.blueTotalExperience, Error.redTotalExperience, c=Error.blueWins)
plt.xlabel("blue experience")
plt.ylabel("red experience")
```

```
Out[477]: Text(0, 0.5, 'red experience')
```




In [478]: Error_normal

Out[478]:

	gameId	blueWins	blueWardsPlaced	blueWardsDestroyed	blueFirstBlood	blueKills	blueKills
4433	4516589637	1	14	1	0	9	
1249	4517512393	1	14	4	1	9	
3860	4522690291	1	12	1	1	7	
7547	4486105297	1	18	2	1	7	
6729	4477434796	0	17	3	0	8	
...
7022	4519668670	0	136	4	1	9	
6334	4487018877	1	16	4	1	5	
905	4516044973	0	13	0	0	5	
9241	4505233035	1	39	0	1	5	
1457	4456960285	1	19	5	0	10	

799 rows × 40 columns



```
In [479]: print("Early JGDiff Avg: %0.4f"
           % (data.blueTotalJungleMinionsKilled-data.redTotalJungleMinionsKilled).mean())
print("Early JGDiff Avg for the holdout: %0.4f"
      % (data_normal.blueTotalJungleMinionsKilled-data_normal.redTotalJungleMinionsKilled).
print("Early JGDiff Avg for Wrong Prediction: %0.4f"
      % (Error_normal.blueTotalJungleMinionsKilled-Error_normal.redTotalJungleMinionsKilled)
print("Early JGDiff Avg for Correct Prediction: %0.4f"
      % (Correct_normal.blueTotalJungleMinionsKilled-Correct_normal.redTotalJungleMinionsKi
```

Early JGDiff Avg: -0.8034

Early JGDiff Avg for the holdout: -1.0067

Early JGDiff Avg for Wrong Prediction: -0.9837

Early JGDiff Avg for Correct Prediction: -1.0152

In []: